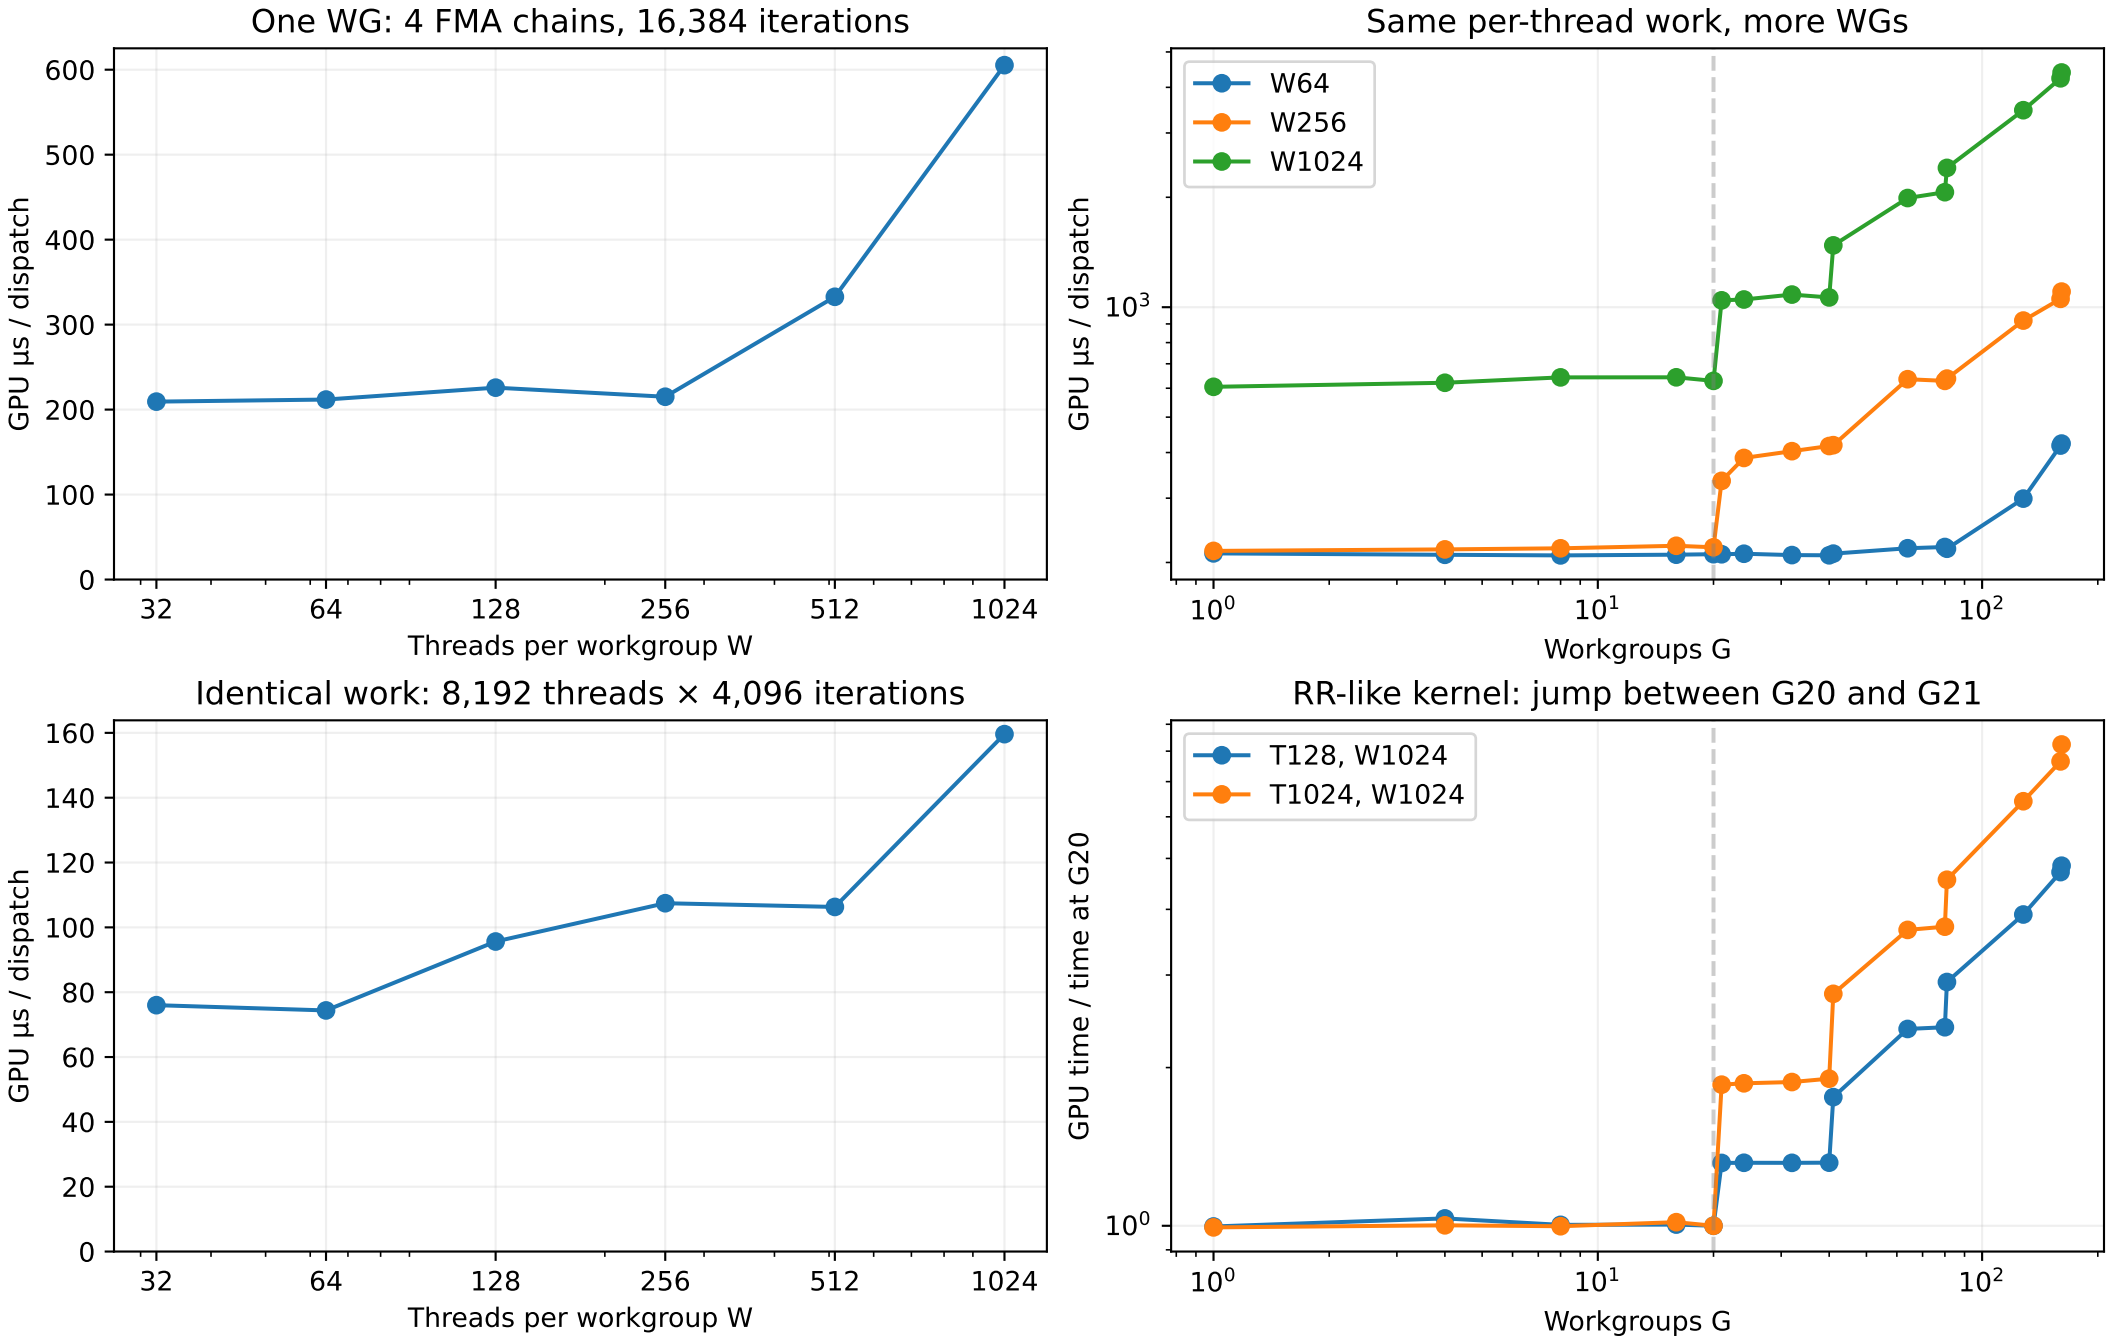
One WG: 4 FMA chains, 16,384 iterations
GPU µs / dispatch
Threads per workgroup W
0
100
200
300
400
500
600
32
64
128
256
512
1024
Same per-thread work, more WGs
GPU µs / dispatch
Workgroups G
W64
W256
W1024
10^3
10^0
10^1
10^2
Identical work: 8,192 threads × 4,096 iterations
GPU µs / dispatch
Threads per workgroup W
0
20
40
60
80
100
120
140
160
32
64
128
256
512
1024
RR-like kernel: jump between G20 and G21
GPU time / time at G20
Workgroups G
T128, W1024
T1024, W1024
10^0
10^0
10^1
10^2