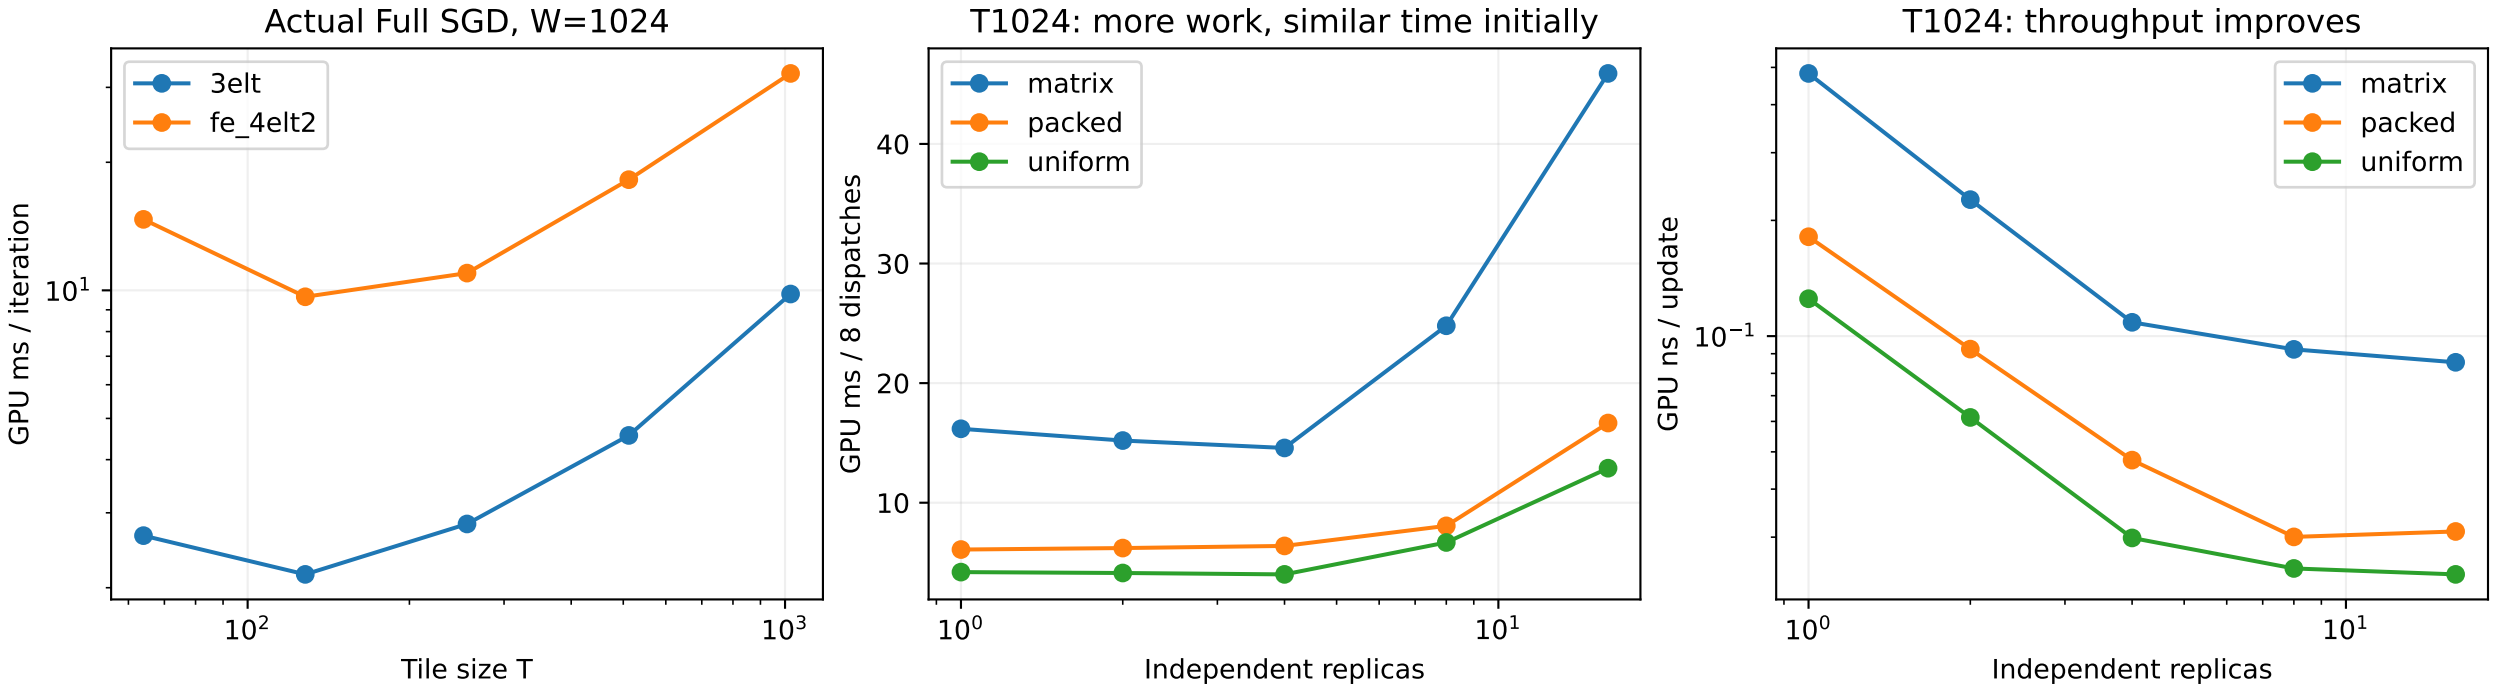
Actual Full SGD, W=1024
3elt
fe_4elt2
GPU ms / iteration
10^1
10^2
10^3
Tile size T
T1024: more work, similar time initially
matrix
packed
uniform
GPU ms / 8 dispatches
40
30
20
10
10^0
10^1
Independent replicas
T1024: throughput improves
matrix
packed
uniform
GPU ns / update
10^{-1}
10^0
10^1
Independent replicas